

# 开展编译器性能分析和优化需要做哪些准备

李永泰

中国科学院软件研究所 PLCT 实验室

[liyongtai@iscas.ac.cn](mailto:liyongtai@iscas.ac.cn)

2025/3/22

# 目录

- 编译环境的搭建
- SPEC CPU 测试套件的编译
- 测试运行
- perf 和 objdump 工具介绍
- 两个例子

# 一些背景

- RISC-V
- Linux
- GCC
- LLVM
- SPEC CPU 2017/2006
- Code size / Dynamic instruction count

# 编译环境的搭建

- 硬件
- 通过包管理安装 gcc/clang 等工具
- 编译 GNU 工具链
- 编译 LLVM 工具链

```
debian@revyos-pioneer:~$ uname -a
Linux revyos-pioneer 6.1.72-pioneer #2024.01.19.06.06+4005122ed SMP Fri Jan 19 07:04:20 UTC 2024 riscv64 GNU/Linux
debian@revyos-pioneer:~$ fastfetch
```

```
_,met$$$$$gg.
,g$$$$$$$$$$$$$P.
,g$$P"          ""Y$$.".
,$$P'           `$$$
',$$P           ,ggs.   `$$b:
`d$$'           ,P""'   $$$
$$P             d$'     $$$P
$$:             $. -    ,d$$'
$$;             Y$b._   _,d$P'
Y$$             `."Y$$$$$P"
`$$b            "-._--
`Y$$
`Y$$
`$$b.
`Y$$b.
`"Y$b._
`""""
```

debian@revyos-pioneer

-----

**OS:** Debian GNU/Linux bookworm 12.0 riscv64  
**Host:** Sophgo Mango  
**Kernel:** 6.1.72-pioneer  
**Uptime:** 7 days, 10 hours, 33 mins  
**Packages:** 1333 (dpkg)  
**Shell:** bash 5.2.15  
**Cursor:** Adwaita  
**Terminal:** /dev/pts/0  
**CPU:** rv64gc (64)  
**GPU:** AMD Radeon HD 6450/7450/8450 / R5 230 OEM  
**Memory:** 715.00 MiB / 124.88 GiB (1%)  
**Swap:** Disabled  
**Disk (/):** 370.38 GiB / 937.56 GiB (40%) - ext4  
**Local IP (enP2p129s0f1):** 192.168.16.119/24 \*  
**Locale:** en\_GB.UTF-8



```
debian@revyos-pioneer:~$
```

```
debian@revyos-pioneer:~$ gcc --version
gcc (Debian 13.2.0-4revyos1) 13.2.0
Copyright (C) 2023 Free Software Foundation, Inc.
This is free software; see the source for copying conditions. There is NO
warranty; not even for MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE.
```

```
debian@revyos-pioneer:~$ clang --version
Debian clang version 14.0.6
Target: riscv64-unknown-linux-gnu
Thread model: posix
InstalledDir: /usr/bin
```

```
debian@revyos-pioneer:~$ cmake --version
cmake version 3.25.1
```

CMake suite maintained and supported by Kitware ([kitware.com/cmake](https://kitware.com/cmake)).

```
debian@revyos-pioneer:~$ ninja --version
1.11.1
```

```
debian@revyos-pioneer:~$ █
```

# 构建 GCC

依赖的基本工具：gcc, g++, make, binutils, libtool, bzip2, gzip, tar, perl

依赖的一些库：GMP、MPFR、MPC

- 使用系统包管理工具安装：libgmp-dev libmpfr-dev libmpc-dev (debian)
- gcc 源码中提供的下载工具

\*<https://gcc.gnu.org/wiki/InstallingGCC>

\*<https://gcc.gnu.org/install/prerequisites.html>

# 构建 GCC

## 获取源码

下载源码包

<https://gcc.gnu.org/mirrors.html>

<https://ftp.gnu.org/gnu/gcc/gcc-14.2.0/>

或者从 git 获取

```
git clone https://github.com/gcc-mirror/gcc  
# git://gcc.gnu.org/git/gcc.git  
# checkout 到感兴趣的 commit/branch
```

下载依赖库

```
./contrib/download_prerequisites
```



# 构建 GCC

- 创建独立构建目录

防止污染源码，编译不同版本时需要清理/创建多个 build 目录

```
mkdir ../build && cd ../build
```

- 运行 configure

```
../gcc/configure --prefix=/path/to/install -enable-languages=c,c++,fortran
```

有非常多的配置选项，可以使用 `configure --help` 查看

- 编译安装

```
make -j$(nproc)  
make install
```

# 构建 LLVM

- 构建 LLVM 的依赖相对较少，准备好 c/c++ 编译器,cmake,ninja 即可
- 获取源码

```
git clone https://github.com/llvm/llvm-project.git  
# https://github.com/llvm/llvm-project/releases/
```

# 构建 LLVM

- 创建独立构建目录并运行 cmake

```
mkdir build && cd build
```

```
cmake -DLLVM_TARGETS_TO_BUILD="RISCV" -DLLVM_ENABLE_PROJECTS="clang;flang;mllr" \
-DCMAKE_BUILD_TYPE=Release -DCMAKE_C_COMPILER=clang -DCMAKE_CXX_COMPILER=clang++ \
-DCMAKE_INSTALL_PREFIX=/path/to/install -G Ninja \
-DLLVM_PARALLEL_LINK_JOBS=N -DLLVM_USE_LINKER=mold ../llvm
```

- 编译安装

```
ninja -j$(nproc)
ninja install
```

编译 LLVM 链接时所需内存较大，尤其是 Debug 版本，需要控制并行链接数

# SPEC CPU 测试套件的编译

## SPEC CPU 概览

- 测试项目分成两类:
  - INT (Integer Performance): 编译、压缩, 逻辑运算.
  - FP (Floating-Point Performance): 科学计算, 图像处理.
- 两个模式:
  - SPECspeed: Single-threaded optimization.
  - SPECrate: Multi-core/thread performance.

# Int 类型测试程序

SPECrate2017 Integer	SPECspeed2017 Integer	Language	KLOC	Application Area
500.perlbench_r	600.perlbench_s	C	362	Perl interpreter
502.gcc_r	602.gcc_s	C	1,304	GNU C compiler
505.mcf_r	605.mcf_s	C	3	Route planning
520.omnetpp_r	620.omnetpp_s	C++	134	Discrete Event simulation - computer network
523.xalancbmk_r	623.xalancbmk_s	C++	520	XML to HTML conversion via XSLT
525.x264_r	625.x264_s	C	96	Video compression
531.deepsjeng_r	631.deepsjeng_s	C++	10	Artificial Intelligence: alpha-beta tree search (Chess)
541.leela_r	641.leela_s	C++	21	Artificial Intelligence: Monte Carlo tree search (Go)
548.exchange2_r	648.exchange2_s	Fortran	1	Artificial Intelligence: recursive solution generator (Sudoku)
557.xz_r	657.xz_s	C	33	General data compression

# fp 类型测试程序

SPECrate2017 Floating Point	SPECspeed2017 Floating Point	Language	KLOC	Application Area
503.bwaves_r	603.bwaves_s	Fortran	1	Explosion modeling
507.cactuBSSN_r	607.cactuBSSN_s	C++, C, Fortran	257	Physics: relativity
508.namd_r		C++	8	Molecular dynamics
510.parest_r		C++	427	Biomedical imaging: optical tomography with finite elements
511.povray_r		C++, C	170	Ray tracing
519.lbm_r	619.lbm_s	C	1	Fluid dynamics
521.wrf_r	621.wrf_s	Fortran, C	991	Weather forecasting
526.blender_r		C++, C	1,577	3D rendering and animation
527.cam4_r	627.cam4_s	Fortran, C	407	Atmosphere modeling
	628.pop2_s	Fortran, C	338	Wide-scale ocean modeling (climate level)
538.imagick_r	638.imagick_s	C	259	Image manipulation
544.nab_r	644.nab_s	C	24	Molecular dynamics
549.fotonik3d_r	649.fotonik3d_s	Fortran	14	Computational Electromagnetics
554.roms_r	654.roms_s	Fortran	210	Regional ocean modeling

# 两种测试模式

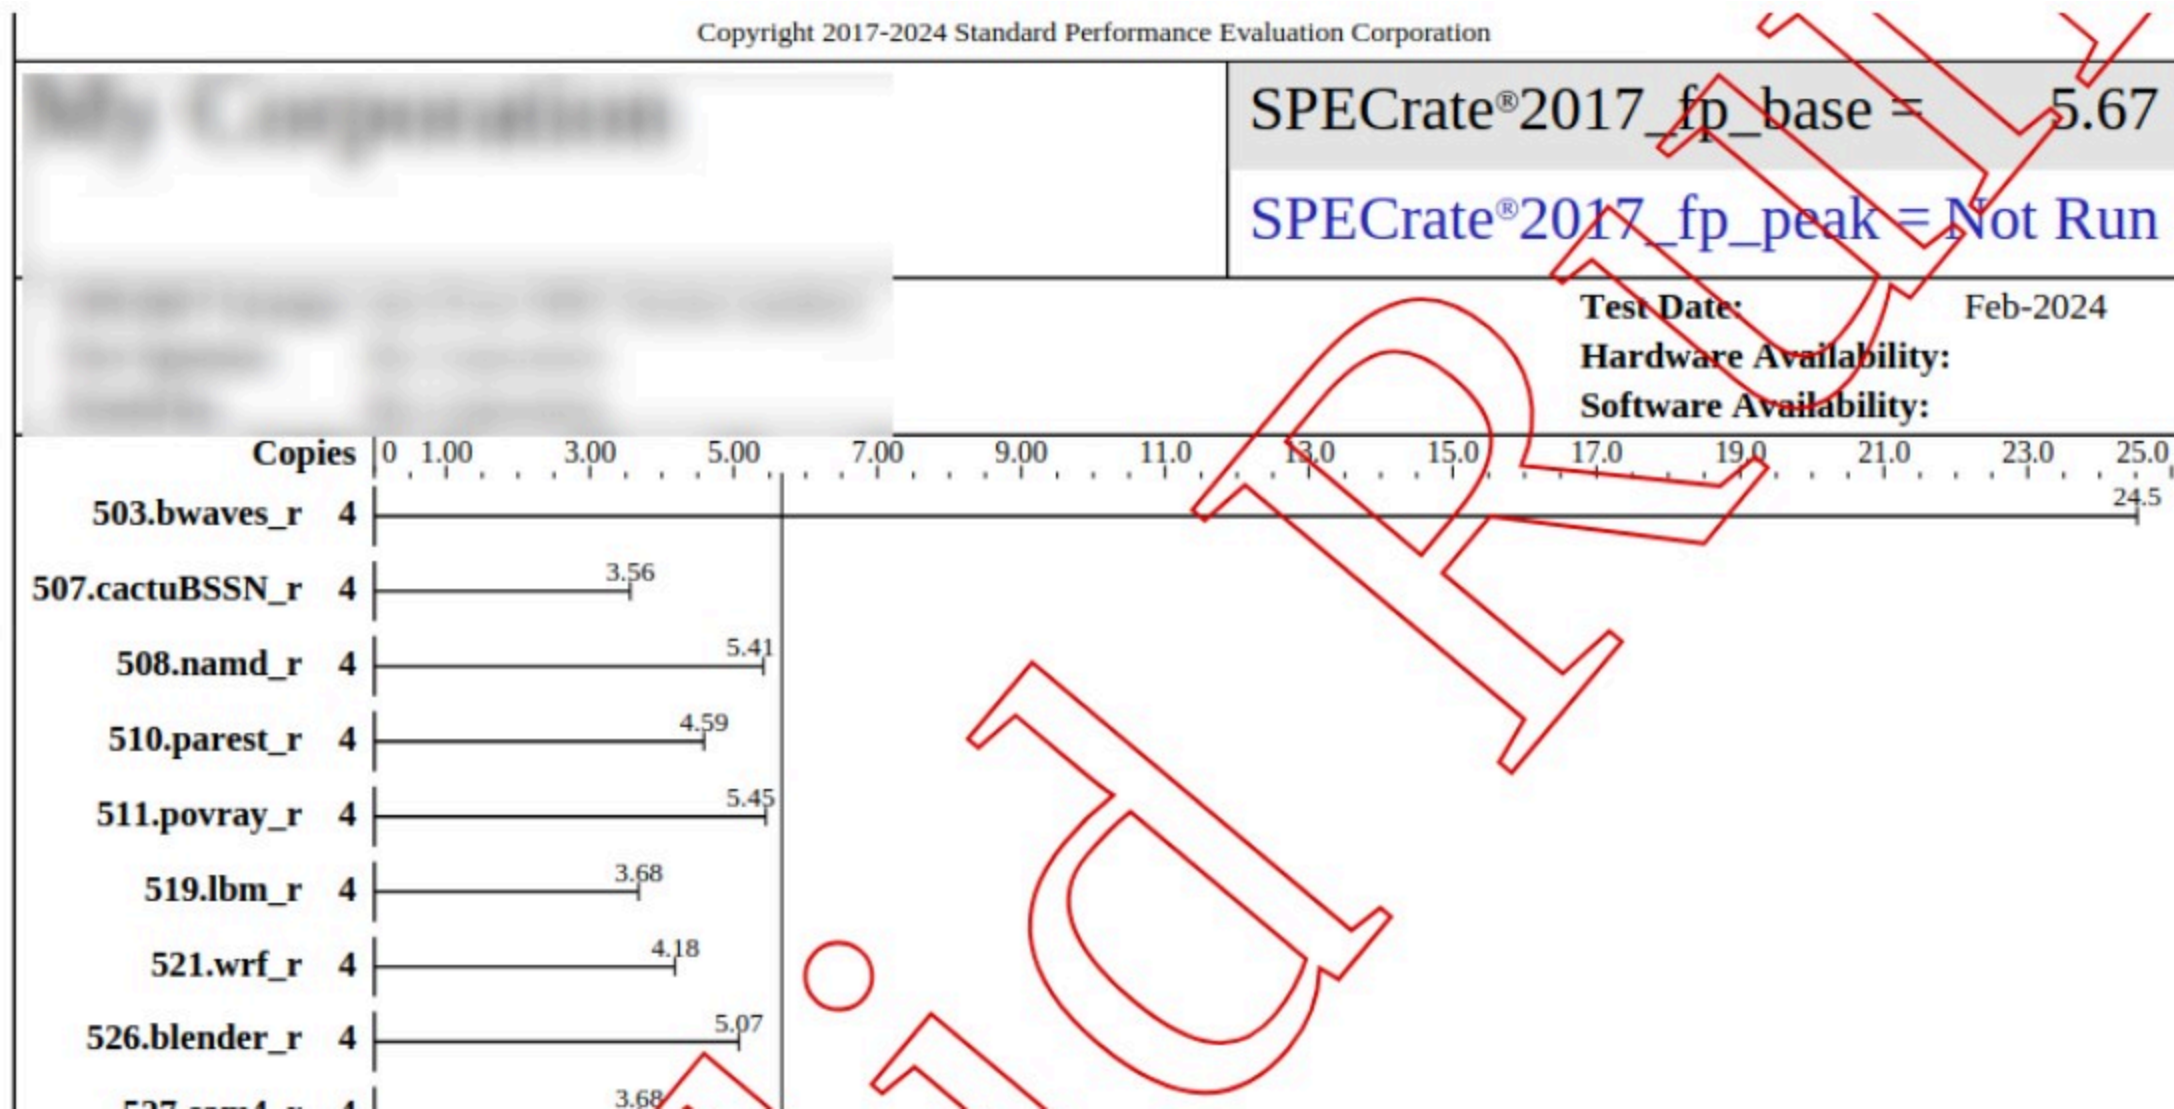
## SPECspeed Metrics

- 1 copy of each benchmark in a suite is run.
- The tester may choose how many OpenMP threads to use.
- For each benchmark, a performance ratio is calculated as:  
`time on a reference machine / time on the SUT`
- Higher scores mean that less time is needed.
- **Example:**
  - The reference machine ran 600.perlbench\_s in 1775 seconds.
  - A particular SUT took about 1/5 the time, scoring about 5.
  - More precisely:  $1775/354.329738 = 5.009458$

## SPECrate Metrics

- The tester chooses how many concurrent copies to run
- OpenMP is disabled.
- For each benchmark, a performance ratio is calculated as:  
`number of copies * (time on a reference machine / time on the SUT)`
- Higher scores mean that more work is done per unit of time.
- **Example:**
  - The reference machine ran 1 copy of 500.perlbench\_r in 1592 seconds.
  - A particular SUT ran 8 copies in about 1/3 the time, scoring about 24.
  - More precisely:  $8*(1592/541.52471) = 23.518776$

# 跑分结果





# SPEC CPU 的安装

- 获取安装包，通常情况下是一个 \*.iso 镜像文件
- 提取文件
- 打 patch 编译工具集（可选）
- 安装、验证

## 哪些情况需要自己编译 RISC-V 工具集

1. SPEC CPU 2006
2. SPEC CPU 2017, version < 1.0.9

低版本可以通过 `runcpu -update` 在线升级到最新版本  
但前提是当前有可用的 `runcpu` 工具

\*<https://www.spec.org/cpu2017/Docs/tools-build.html>

## 提取安装文件和 tools 源码

```
mount /path/to/cpu2017.iso /mnt -o loop
mkdir /home/test/spec2017
cp -r /mnt/* /home/test/spec2017_iso

chmod -R +w /home/test/spec2017_iso #增加写权限

export SPEC=/home/test/spec2017_iso
cd $SPEC/install_archives
tar xf tools-src.tar -C .. #解压工具集源码
```

# 给源码打补丁

1. config.guess 和 config.sub 太过老旧, 不能识别 riscv 架构

下载最新版本的 config.guess 和 config.sub

[http://git.savannah.gnu.org/gitweb/?p=config.git;a=blob\\_plain;f=config.guess](http://git.savannah.gnu.org/gitweb/?p=config.git;a=blob_plain;f=config.guess)

[http://git.savannah.gnu.org/gitweb/?p=config.git;a=blob\\_plain;f=config.sub](http://git.savannah.gnu.org/gitweb/?p=config.git;a=blob_plain;f=config.sub)

替换以下位置

```
$SPEC/tools/src/specinvoke/  
$SPEC/tools/src/specsum/build-aux/  
$SPEC/tools/src/tar-1.28/build-aux/  
$SPEC/tools/src/make-4.2.1/config/  
$SPEC/tools/src/rxp-1.5.0/  
$SPEC/tools/src/expat-2.1.0/conftools/  
$SPEC/tools/src/xz-5.2.2/build-aux/
```

# 给源码打补丁

2. 编译 perl 时, 会将 gcc 10 识别为 gcc 1.x

这是因为 \$SPEC/tools/src/perl-5.24.0/Configure 和 \$SPEC/tools/src/perl-5.24.0/cflags.SH 中使用 1\* 匹配 gcc 版本号, 需改为 1.\*

```
1.1
$rm -f try try.*
case "$gccversion" in
-1*) cpp=`./loc gcc-cpp $cpp $pth` ;;
+1.*) cpp=`./loc gcc-cpp $cpp $pth` ;;
esac
case "$gccversion" in
'' ) gccosandvers='' ;;
@@ -5438,7 +5438,7 @@ fi
case "$hint" in
default|recommended)
    case "$gccversion" in
-    1*) dflt="$dflt -fpcc-struct-return" ;;
+    1.*) dflt="$dflt -fpcc-struct-return" ;;
    esac
    case "$optimize:$DEBUGGING" in
*-g*:old) dflt="$dflt -DDEBUGGING";;
@@ -5453,7 +5453,7 @@ default|recommended)
```

## 给源码打补丁

3. \$SPEC/tools/src/TimeDate-2.30/t/getdate.t 测试失败 需要做以下修改

```
--- getdate.t
+++ getdate.t
@@ -156,7 +156,7 @@ Jul 22 10:00:00 UTC 2002          ;1027332000
!;

require Time::Local;
-my $offset = Time::Local::timegm(0,0,0,1,0,70);
+my $offset = Time::Local::timegm(0,0,0,1,0,1970);

@data = split(/\n/, $data);
```

\*<https://github.com/sihuan/llvm-work/blob/master/spec2017/README.md>

# 编译并打包工具集

1. 编译, 使用 \$SPEC/tools/src/buildtools

```
export MAKEFLAGS=-j4  
./buildtools
```

2. 打包, 使用 \$SPEC/bin/packagetools

```
mkdir -p $SPEC/tools/bin/linux-riscv64  
echo '.....' > $SPEC/tools/bin/linux-riscv64/description  
./bin/packagetools linux-riscv64
```

# 安装 SPEC CPU

```
./install.sh -u linux-riscv64 -d /home/test/spec2017
```

## 测试

```
source shrc
runcpu --version
SPEC CPU(r) 2017 Benchmark Suites
Copyright 1995-2019 Standard Performance Evaluation Corporation (SPEC)

runcpu v6612
Using 'linux-riscv64' tools
    This is the SPEC CPU2017 benchmark tools suite.
```



# 编写成对的配置文件

1. base 和 peak 两种配置
  2. 主要需要修改的地方
    - 编译器路径
    - -march
    - 不同的优化选项、优化等级
- 可以查看 Example-\* 开头的样例配置

```

$ diff llvm-c9a6e993f7b3-rv64gc_zba_zbb_zbs.cfg gcc-d28ea8e5a704-rv64gc_zba_zbb_zbs.cfg
1c1
< %define commit c9a6e993f7b3
---
> %define commit d28ea8e5a704
<-SKIP->
60,65c60,65
<   preENV_LD_LIBRARY_PATH = %{llvm_dir}/lib/:lib
<   SPECLANG                = %{llvm_dir}/bin/
<   CC                      = $(SPECLANG)clang -std=c99 -Wno-implicit-int
<   CXX                     = $(SPECLANG)clang++ -std=c++03
<   FC                      = $(SPECLANG)flang-new
---
>
>   preENV_LD_LIBRARY_PATH = %{gcc_dir}/lib/:lib
>   SPECLANG                = %{gcc_dir}/bin/
>   CC                      = $(SPECLANG)gcc -L%{gcc_dir}/lib -std=c99 -Wno-implicit-int
>   CXX                     = $(SPECLANG)g++ -L%{gcc_dir}/lib -std=c++03
>   FC                      = $(SPECLANG)gfortran -L%{gcc_dir}/lib

```

# 跑个分看看

```
source shrc
ulimit -s unlimited
runcpu --config=xxx --size=ref --tuning=all all
```

测试范围可以形如

- 503
- 519.lbm\_r
- int
- fprate

\*如果想要生成一个合法的报告，需要跑多次（默认），test size 需要是 ref，tuning 可以是 base 或者 all

\*<https://www.spec.org/cpu2017/Docs/runcpu.html>

# 分步运行

- build —— 仅**编译**指定的测试项目
- runsetup —— 构建测试运行环境

## 仅编译测试











使用以下命令编译测试并打包编译得到的可执行程序

```
runcpu --config label.cfg --action build intspeed  
spectar cvf - config/label.cfg benchspec/CPU/*/exe/*label | specxz -T 0 > label.tar.xz
```

上边的 *label* 形如 `llvm-c9a6e993-rv64gc_zba_zbb_zbs` or `gcc-d28ea8e5-rv64gcv_zba_zbb_zbs` .

action 选项:

- build: 构建指定套件
- buildsetup: 生成构建目录, 复制源码, 生成 Makefile

✓-  benchspec	17.3 MiB
✓-  CPU	17.3 MiB
✓-  500.perlbench_r	2.2 MiB
✓-  exe	2.2 MiB
 perlbench_r_base.llvm-c9a6e993f7b3-rv64gc_zba_zbb_zbs-64	2.2 MiB
✓-  502.gcc_r	9.9 MiB
✓-  exe	9.9 MiB
 cpugcc_r_base.llvm-c9a6e993f7b3-rv64gc_zba_zbb_zbs-64	9.9 MiB
>-  503.bwaves_r	1.8 MiB
>-  SKIP	3.4 MiB
✓-  config	60.9 KiB
 llvm-c9a6e993f7b3-rv64gc_zba_zbb_zbs.cfg	60.9 KiB

## 生成运行环境

生成运行环境，包括测试需要的输入和 `speccmds.cmd` 等文件.

```
runcpu --config label.cfg --action runsetup intspeed
```

`--action runsetup` 选项生成运行目录，复制二进制文件和数据文件，创建控制文件

测试数据和控制文件是架构无关的，但是 test size 有关 (ref/test)

## 548\_exchange2\_r 的运行目录

---



compare.cmd



control



exchange2\_r.exe



puzzles.txt



speccmds.cmd

- `exchange2_r.exe` : 待测试的可执行程序
- `puzzles.txt` : 输入数据
- `control` , `speccmds.cmd` , `compare.cmd` : 一些控制文件, 描述了测试如何运行, 如何对比数据



## 生成运行环境

使用 `specinvoke` 可以查看 `speccmds.cmd` 中关于测试如何运行的描述

```
$ specinvoke -n speccmds.cmd
# specinvoke r4356
# Invoked as: specinvoke -n speccmds.cmd
# timer ticks over every 1000 ns
# Use another -n on the command line to see chdir commands and env dump
# Starting run for copy #0
../run_base_refrate_llvm-c9a6e993f7b3-rv64gc_zba_zbb_zbs-64.0000/\
exchange2_r_base.llvm-c9a6e993f7b3-rv64gc_zba_zbb_zbs-64 6 > exchange2.txt 2>> exchange2.err
specinvoke exit: rc=0
```

现在我们已经可以手动去运行任意一个测试。

# 为什么选用 QEMU

1. RISC-V 机器目前还是比较慢
2. 很多测试项有多个 workload，可以设计成并行运行
3. 借助 QEMU 的插件系统可以方便的统计指令数（也可以自己编写插件统计更详细的信息）

## 运行测试并获取数据

- Code Size: strip 可执行文件然后直接统计文件大小
- 动态指令数 —— 使用 QEMU User 运行并统计
  - 构建 QEMU 和 insn 插件
  - 带 insn 插件运行测试

# 构建带有 TCG 插件的 QEMU

## 获取源码

```
git clone https://github.com/qemu/qemu.git && cd qemu
mkdir build && cd build

../configure --prefix=/path/to/install \
    --target-list=riscv64-linux-user --enable-plugins

make
make install
```

\*<https://www.qemu.org/docs/master/devel/index-build.html>

\*<https://www.qemu.org/docs/master/devel/tcg-plugins.html>

\*<https://github.com/qemu/qemu/blob/master/tests/tcg/plugins/insn.c>

## 使用帶有 insn 插件的 QEMU

```
$ path/to/qemu-riscv64 -plugin path/to/plugin/libinsn.so -d plugin ./demo  
cpu 0 insns: 20250322  
total insns: 20250322
```

## 更多的选项

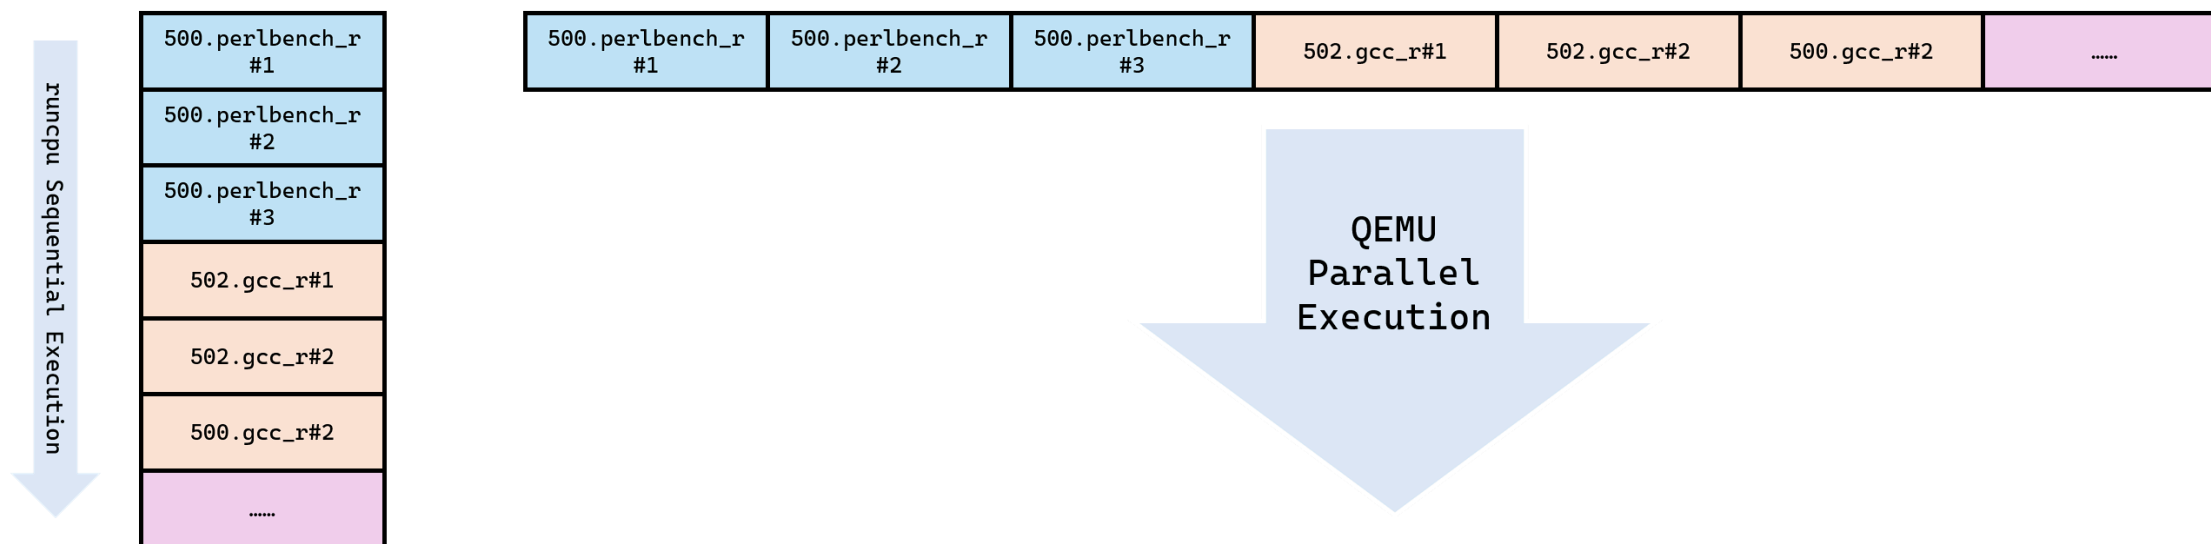
实际运行 SPEC CPU 测试的话，需要更多的选项，比如

- `-L /opt/riscv/sysroot`
- `-cpu rv64,v=true,vext_spec=v1.0`
- `-s 819200000000`

不过可以用自动化工具来代劳

# 开源的自动化工具

<https://github.com/sihuan/countspect>







## 统计得到的数据

**Table 1: LLVM Code Size Relative to GCC (Scalar)**


<b>Range</b>	<b>C/C++</b>	<b>Fortran</b>
$< 0.9$	508,541,507,519,531 557,538,525,505,544	527
0.9-1	500,502	
1-1.1	511,510,520	
$> 1.1$	523,526	521,554,549,548,503

Google表格




Benchmark	gcc s	gcc v	llvm s	llvm v	llvms/gccs	llvmv/gccv
508.namd_r	645960	658328	408984	426632	63.31%	64.81%
541.leela_r	105288	109464	80056	83656	76.04%	76.42%
527.cam4_r	18509560	19053760	11304944	11878504	61.08%	62.34%
519.lbm_r	18864	18952	14808	15088	78.50%	79.61%
531.deepsjeng_r	76120	80296	60992	79288	80.13%	98.74%
557.xz_r	137544	141736	113904	119952	82.81%	84.63%
538.imagick_r	1400704	1429464	1195560	1206960	85.35%	84.43%
525.x264_r	1015032	1076800	883728	936640	87.06%	86.98%
505.mcf_r	22808	22896	19888	21144	87.20%	92.35%
544.nab_r	91712	91808	82272	83744	89.71%	91.22%
500.perlbench_r	2276592	2309456	2062088	2075928	90.58%	89.89%
502.gcc_r	9562056	9635904	9071992	9368568	94.87%	97.23%
511.povray_r	1013680	1038416	1044176	1075616	103.01%	103.58%
510.parest_r	1514696	1599968	1625808	1740664	107.34%	108.79%
520.omnetpp_r	1633032	1649488	1767824	1774944	108.25%	107.61%
523.xalancbmk_r	3156704	3226480	3667744	3726384	116.19%	115.49%
526.blender_r	14898976	15154768	17502152	17710776	117.47%	116.87%
Benchmark	gcc s	gcc v	llvm s	llvm v	llvms/gccs	llvmv/gccv
527.cam4_r	18509560	19053760	11304944	11878504	61.08%	62.34%
521.wrf_r	30947416	35897608	41985072	48908776	135.67%	136.25%
554.roms_r	840072	1311384	2109832	2393064	251.15%	182.48%
549.fotonik3d_r	299200	326704	891232	957944	297.87%	293.21%
548.exchange2_r	127456	147800	1428760	1464288	1120.98%	990.72%
503.bwaves_r	27136	44360	549880	560992	2026.39%	1264.63%

# 动态指令数



# 动态指令数



## 动态指令数

**Table 2: V-Ext DIC Reduction**

<b>Suit</b>	<b>GCC</b>	<b>LLVM</b>
<i>int-rate avg.</i>	6.45%	4.22%
<i>fp-rate avg.</i>	11.73%	16.84%
<i>all avg.</i>	9.43%	11.35%

# 性能分析工具

- perf —— 动态分析，指令数、分支数、热点函数
- objdump —— 静态分析，反汇编，局部分析

# perf

- Linux 内核自带的性能分析工具
- 基于硬件性能计数器、跟踪点 (tracepoints) 进行采样和分析

核心功能:

- CPU 性能分析 (指令、缓存、分支预测)
- 函数级代码热点分析
- 系统级资源监控

# 安装和配置

不同发行版打包情况不同

```
apt install linux-perf # revyos
```

```
pacman -S perf # archlinux
```

配置权限

```
# 允许非 root 用户使用  
sudo sysctl -w kernel.perf_event_paranoid=1
```



# 常用命令

- perf list 列出支持的事件
- perf stat 统计程序运行时的事件
- perf record 采样并保存数据（生成 perf.data）
- perf report 分析采样数据
- perf top 实时显示系统热点

## 其他选项

- perf stat -e 选项可以指定关心的事件
- perf record -F 选项可以控制采样频率

# objdump

GNU Binutils 工具集的一部分，用于解析二进制文件（可执行文件、目标文件、共享库）。

支持反汇编、查看符号、调试信息等等

和 perf 是一个互补的关系，通过 perf 定位运行时的性能问题，然后通过 objdump 反汇编分析代码静态结构，验证编译器优化效果。

# 安装和基础用法

## 安装

```
apt install binutils #一般已经预装/ debian  
pacman -S riscv64-linux-gnu-binutils #x86 Archlinux
```

## 基础用法

```
objdump [选项] <二进制文件>
```

## 常见选项

-d, --disassemble	反汇编代码段
-S, --source	同时显示源码与汇编
-t, --syms	查看符号表
-j, --section	分析指定的节（如 .text, .data）

## demo 程序分析

```
static inline void bar() {  
    return;  
}  
  
int main() {  
    int i;  
    for(i = 0; i < 100000000; i++)  
        bar();  
    return 0;  
}
```

1. 指令数 16xxxxxxxxxxx 是怎么来的
2. 反汇编指定函数 bar
3. 通过符号表检查内联效果

## 实际案例分析 548/473

548.exchange2\_r 是一个针对 9\*9 数独的解题程序，使用 Fortran95 编写，约有 1600 行代码，代码严重依赖递归（最高达 8 层），另外该程序不包括任何浮点运算。

无论是否开启 v 拓展 LLVM 和 GCC 的差异都巨大。

*实际上，这个问题和 B 拓展也没关系，甚至和 RISC-V 架构无关*

[https://www.spec.org/cpu2017/Docs/benchmarks/548.exchange2\\_r.h](https://www.spec.org/cpu2017/Docs/benchmarks/548.exchange2_r.h)

# 首先要解决的一些问题

- 使用 perf 分析的话需要在物理机器上, 但是 SG2042 并没有 B 拓展
  - 尝试在 rv64gc 上复现.
- 手动编译流程(--action buildsetup / 查看 make.out)

```
flang-new -c -o exchange2.fppized.o -march=rv64gc -Ofast exchange2.fppized.f90  
flang-new -march=rv64gc -Ofast exchange2.fppized.o -o exchange2_r
```

- 手动测试流程(specinvoke 读取)

```
./exchange2_r 0 # test size, solve the first problem in `puzzles.txt`  
./exchange2_r 6 # ref size, solve all problems in `puzzles.txt`
```

用 `perf` 在 test size 下录得的数据:

```
perf stat ./exchange2_r 0
perf report
```

GCC:

#	Overhead	Command	Shared Object	Symbol
	38.96%	exchange2_r	exchange2_r	[.] __brute_force_MOD_digits_2.constprop.4.isra.0
	19.95%	exchange2_r	exchange2_r	[.] __brute_force_MOD_digits_2.constprop.3.isra.0
	9.03%	exchange2_r	exchange2_r	[.] __brute_force_MOD_digits_2.constprop.6.isra.0
	7.10%	exchange2_r	libgfortran.so.5.0.0	[.] gfortran_mminloc0_4_i4
	6.25%	exchange2_r	exchange2_r	[.] __logic_MOD_new_solver
	5.50%	exchange2_r	exchange2_r	[.] __brute_force_MOD_digits_2.constprop.5.isra.0
	4.41%	exchange2_r	exchange2_r	[.] specific.4
	2.15%	exchange2_r	exchange2_r	[.] __brute_force_MOD_digits_2.constprop.7.isra.0
	2.07%	exchange2_r	exchange2_r	[.]
		\$xrv64i2p1_m2p0_a2p1_f2p2_d2p2_c2p0_zicsr2p0_zifencei2p0_zmmullp0		
	1.21%	exchange2_r	exchange2_r	[.] hidden_pairs.2
	0.85%	exchange2_r	exchange2_r	[.] naked_triplets.1
	0.71%	exchange2_r	exchange2_r	[.] __brute_force_MOD_brute
	0.64%	exchange2_r	exchange2_r	[.] __brute_force_MOD_digits_2.constprop.2.isra.0
	0.33%	exchange2_r	exchange2_r	[.] __brute_force_MOD_digits_2.constprop.1.isra.0
	0.19%	exchange2_r	exchange2_r	[.] __brute_force_MOD_covered.constprop.0.isra.0
	0.12%	exchange2_r	exchange2_r	[.] __brute_force_MOD_rearrange.isra.0

LLVM:

#	Overhead	Command	Shared Object	Symbol
	88.78%	exchange2_r	exchange2_r	[.] _QMbrute_forcePdigits_2
	5.78%	exchange2_r	exchange2_r	[.] _QMlogicFnew_solverPspecific
	1.34%	exchange2_r	exchange2_r	[.] _QMlogicFnew_solver
	0.67%	exchange2_r	exchange2_r	[.] _QMlogicFnew_solverPhidden_triplets
	0.33%	exchange2_r	exchange2_r	[.] _QMlogicFnew_solverPhidden_pairs
	0.28%	exchange2_r	exchange2_r	[.] _QMlogicFnew_solverPnaked_triplets
	0.28%	exchange2_r	exchange2_r	[.] _QMlogicFnew_solverPnaked_pairs
	0.20%	exchange2_r	exchange2_r	[.] Fortran::runtime::Assign
	0.19%	exchange2_r	exchange2_r	[.] _QMbrute_forcePbrute
	0.16%	exchange2_r	exchange2_r	[.] Fortran::runtime::ReduceDimToScalar<int,
		Fortran::runtime::ExtremumLocAccumulator<Fortran::runtime::NumericCompare<int, true, false> > >		
	0.15%	exchange2_r	libc.so.6	[.] memcpy
	0.15%	exchange2_r	libc.so.6	[.] memset
	0.15%	exchange2_r	exchange2_r	[.] _QMlogicFnew_solverPx_wing
	0.12%	exchange2_r	libc.so.6	[.] _int_free
	0.11%	exchange2_r	libc.so.6	[.] malloc
	0.10%	exchange2_r	exchange2_r	[.] _QMbrute_forcePcovered

GCC:

```
$ sudo perf stat -B -e cache-references,cache-misses,cycles,instructions,branches,faults,migrations ./exchange2_r 6
Performance counter stats for './exchange2_r 6':
    417,772,422      cache-references
    4,425,518        cache-misses          #    1.059 % of all cache refs
  1,327,980,494,790  cycles
  2,081,545,960,539  instructions          #    1.57  insn per cycle
    274,477,603,251  branches
         90          faults
         0           migrations
    664.238029637    seconds time elapsed
    664.003383000    seconds user
     0.015994000    seconds sys
```

LLVM:

```
$ sudo perf stat -B -e cache-references,cache-misses,cycles,instructions,branches,faults,migrations ./exchange2_r 6
Performance counter stats for './exchange2_r 6':
    656,853,453      cache-references
    6,735,733        cache-misses          #    1.025 % of all cache refs
  2,373,045,060,474  cycles
  4,328,214,832,776  instructions          #    1.82  insn per cycle
    496,851,214,471  branches
         83          faults
         0           migrations
   1186.975610235    seconds time elapsed
   1186.570645000    seconds user
     0.007997000    seconds sys
```



# 反汇编分析

GCC:

```
$ objdump --disassemble=__brute_force_MOD_digits_2.isra.0 exchange2_r | wc -l
2312
$ for i in {1..7}; do objdump --disassemble=__brute_force_MOD_digits_2.constprop.${i}.isra.0 exchange2_r | wc -l; done
1094
922
1094
1145
752
925
1025
```

LLVM:

```
$ objdump --disassemble=_QMbrute_forcePdigits_2 exchange2_r | wc -l
2794
```

## 推测原因

在 GCC 上, 热点函数 `digits_2` 被拆成了几个特化版本. 这种拆分特化是过程间常量传播优化造成的 (IPA-CP)。此优化的主要作用之一就是条件分支消除。

也正是这样，特化的函数静态汇编行数也较少。

LLVM 里对应的优化是 `IPSCCP` Pass.

赞美 LLM

# 验证推测

通过使用 `-fno-ipa-cp` 选项关闭 gcc 优化

	gcc -fno-ipa-cp	gcc
exchange2_r 0	93,554,141,493	55,981,214,885

## 尝试在 LLVM 里重复优化

```
flang-new -c -emit-llvm -o exchange2.fppized.ll -march=rv64gc -Ofast exchange2.fppized.f90  
opt -passes="ipsccp" exchange2.fppized.ll -o exchange2.fppized.ipsccp.ll  
flang-new -march=rv64gc -Ofast fppized.ipsccp.ll -o exchange2_r
```

	llvm	llvm + ipsccp
exchange2_r 0	114,450,486,604	70,380,347,586

指令数减少且通过反汇编可以看到 `digits_2` 被特化成形如 `_QMbrute_forcePdigits_2.specialized.3` 的几个函数。

```
llvm-project / llvm / lib / Passes / PassBuilderPipelines.cpp

Code Blame 2189 lines (1824 loc) · 90.9 KB

402 PassBuilder::build01FunctionSimplificationPipeline(OptimizationLevel Level,
1110 // tests in icf sequences.
1116 if (Phase == ThinOrFullLTOPhase::ThinLTOPostLink)
1117     MPM.addPass(LowerTypeTestsPass(nullptr, nullptr, true));
1118
1119 invokePipelineEarlySimplificationEPCallbacks(MPM, Level);
1120
1121 // Interprocedural constant propagation now that basic cleanup has occurred
1122 // and prior to optimizing globals.
1123 // FIXME: This position in the pipeline hasn't been carefully considered in
1124 // years, it should be re-analyzed.
1125 MPM.addPass(IPSCCPass(
1126     IPSCCOptions(/*AllowFuncSpec=*/|
1127                 Level != OptimizationLevel::Os &&
1128                 Level != OptimizationLevel::Oz &&
1129                 !isLTOPreLink(Phase))););
1130
1131 // Attach metadata to indirect call sites indicating the set of functions
1132 // they may target at run-time. This should follow IPSCCP.
1133 MPM.addPass(CalledValuePropagationPass());
1134
```

```
6 ■■■■■■ llvm/lib/Passes/PassBuilderPipelines.cpp

1204 1204 else
1205 1205     MPM.addPass(buildInlinerPipeline(Level, Phase));
1206 1206
1207 + MPM.addPass(IPSCCPass(
1208 +     IPSCCOptions(/*AllowFuncSpec=*/
1209 +                 Level != OptimizationLevel::Os &&
1210 +                 Level != OptimizationLevel::Oz &&
1211 +                 !isLTOPreLink(Phase))););
1212 +
1207 1213 // Remove any dead arguments exposed by cleanups, constant folding globals,
1208 1214 // and argument promotion.
1209 1215 MPM.addPass(DeadArgumentEliminationPass());
```

IPSCCP Pass 默认在较早的位置启用，尝试重复运行 Pass 可以得到一个比较好的效果。

```
$ objdump -D exchange2_r_patched_llvm | grep "digits_2.*:$"  
00000000000011ab0 <_QMbrute_forcePdigits_2>:  
00000000000018a4e <_QMbrute_forcePdigits_2.specialized.1>:  
00000000000019820 <_QMbrute_forcePdigits_2.specialized.2>:  
0000000000001a436 <_QMbrute_forcePdigits_2.specialized.3>:  
0000000000001ae78 <_QMbrute_forcePdigits_2.specialized.4>:  
0000000000001ba8e <_QMbrute_forcePdigits_2.specialized.5>:  
0000000000001c7e6 <_QMbrute_forcePdigits_2.specialized.6>:  
0000000000001d072 <_QMbrute_forcePdigits_2.specialized.7>:  
0000000000001dad0 <_QMbrute_forcePdigits_2.specialized.8>:
```

workaround 效果：和 GCC 表型类似，性能提升 ~ 40%

重新用 perf 在 rv64gc 上  
统计的 exchange2\_r 0 数  
据  
额外的 x86\_64 数据

Compiler	Instructions on rv64gc
GCC #d28ea8e5	55,965,728,914
LLVM #62d44fbd	105,416,890,241
LLVM #62d44fbd with patch	62,693,427,761

Compiler	cpu_atom instructions on x86_64
LLVM #62d44fbd	100,147,914,793
LLVM #62d44fbd with patch	53,077,337,115

## 473 分析

1. perf stat 查看动态指令数对比
2. perf record/report 定位问题函数
3. objdump 局部静态分析
4. 通过汇编上的差异 推断 关联的优化技术（可以请 DeepSeek 帮忙）
5. 阅读源码，编写较小的复现程序



# 总结和拓展

1. 构建待测的编译器
2. 构建测试套件
3. 运行并统计指标信息
4. 对比数据，使用 perf 动态分析，缩小范围
5. 反汇编，分析代码，找到汇编层面可能造成性能差异的点
6. 反向查找对应的优化技术
7. 构建复现样例

Code Size 数据: <https://docs.google.com/spreadsheets/d/1e6sAkT1kZa8LQo4MWgT-NomF8fSHnClrJMVTrxktUAM>

动态指令数数据:

[https://docs.google.com/spreadsheets/d/1BSSc5XRr\\_QUmEgupRs3MgUJ4pICWsNW\\_X25vADO7DBY](https://docs.google.com/spreadsheets/d/1BSSc5XRr_QUmEgupRs3MgUJ4pICWsNW_X25vADO7DBY)

countspec: <https://github.com/sihuan/countspec>

Workaround for 548: <https://github.com/llvm/llvm-project/pull/96620>

# Thanks